

Chapter 2

End to End Machine Learning

DuLi

2019 年 3 月 25 日

目录	2
----	---

目录

1 Outline	3
2 Working with Real Data	3
3 Look at the Big Picture	4
3.1 Frame the Problem	4
3.2 Select a Performance Measure	5
4 Get the data	6
4.1 Create the workspace	6
4.2 Take a Quick Look at the Data Structure	6
4.3 Create a Testset	9

1 Outline

Here are main steps you will go through:

1. Look at the big picture.
2. Get the data.
3. Discover and visualize the data to gain insights.
4. Prepare the data for machine learning algorithms.
5. Select a model and train it.
6. Fine-tune your model.
7. Present your solution.
8. Launch, monitor, and maintain your system.

2 Working with Real Data

Here are a few places you can look to get data:

- Popular open data open repositories:
 - UC Irvine Machine Learning Repositories.
 - Kaggle datasets.
 - Amazon's AWS datasets.
- Meta Portals (they list open data repositories)
 - dataportals.org
 - opendatamonitor.eu
 - quandl.com
- Other pages listing many popular open data repositories

- Wikipedia's list of Machine Learning datasets.
- Quora.com question.
- Datasets subreddit.




图 1: California Housing Prices Databases

3 Look at the Big Picture

3.1 Frame the Problem

The first question to ask you is what exactly is the business objective; building a model is probably not the end goal. How do you expect use and benefit from this model? This is important because it will determine how you frame the problem, what algorithms you will select, what performance measure you will use to evaluate your model, and how much effort you should spend tweaking it. Your model output (a predicting of a district's median housing price) will be fed to another Machine Learning system along with

many other signals. This downstream system will determine whether it is worth investing in a given area or not.

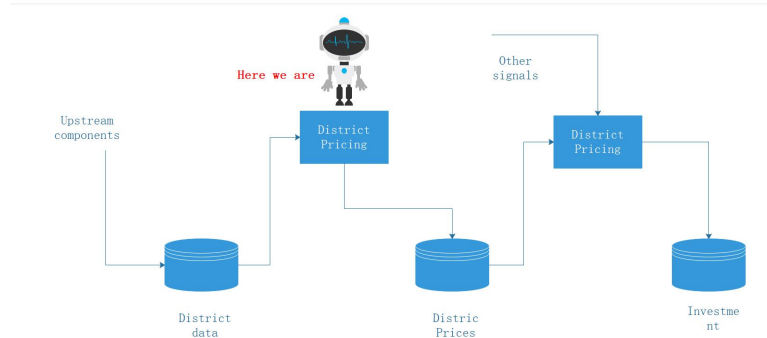


图 2: A machine learning pipeline for real estate investment

The next step is framing the problem: is it supervised, unsupervised, or Reinforcement learning? Is it a classification task, a regression task, or something else? Should we use batch learning or online learning technique? Clearly it is supervised, we need historical data to train the model. Moreover it is a typical regression task, more specifically, this is a multivariate regression since the system will use multiple features to make a prediction. In the first chapter, we predicted life satisfaction based on just one feature, the GDP per capita. Finally, there is no continuous flow of data coming in the system, there is no particular need to adjust to changing data rapidly, and the data is small enough to fit in memory, so plain batch learning should do just fine.

3.2 Select a Performance Measure

The next step is to select a performance measure. A typical performance measure for regression problems is the Root Mean Square Error (均方根误差), it measures the standard deviation of the errors the system makes in its prediction.

$$RMSE(X, h) = \sqrt{\frac{1}{m} \sum_{i=1}^m (h(X^i) - y^i)^2} \quad (1)$$

Even though the RMSE is generally the preferred performance measure for regression tasks, in some

contexts you may prefer to use another function. For example, suppose that there are many outlier districts. In that case, we may consider using the Mean Absolute Error (平均绝对误差):

$$MAE(X, h) = \frac{1}{m} \sum_{i=1}^m |h(X^i) - y^i| \quad (2)$$

Both the RMSE and MAE are ways to measure the *distance* between two vectors. Various distance measures or norms are possible:

- Euclidean norm (欧几里得距离)。
- Manhattan norm (曼哈顿距离), it measures the distance between two points in a city if you can only travel along orthogonal city blocks. 也就是指城市中两点之间沿着街区边缘走路的距离。
- More generally, the l_k norm of a vector v containing n elements is defined as:

$$\|v\| = (|v_0|^k + |v_1|^k + \cdots + |v_n|^k)^{\frac{1}{k}} \quad (3)$$

4 Get the data

It's time to get your hands dirty.

4.1 Create the workspace

First, you need to have Python environment installed. We recommend you installed anaconda on your computer. <https://www.anaconda.com/>

4.2 Take a Quick Look at the Data Structure

Let's take a look at the Data Structure. I download the database from Kaggle. <https://www.kaggle.com/camnugent/california-housing-prices>. Let's take a glance of the data structure. Each row represents one district. There are 10 attributes (Figure 3): longitude, latitude, housing_median_age, total_rooms, total_bedrooms, population, households, median_income, median_house_value, ocean_proximity.

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0	NEAR BAY
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0	NEAR BAY
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0	NEAR BAY
5	-122.25	37.85	52.0	919.0	213.0	413.0	193.0	4.0368	269700.0	NEAR BAY
6	-122.25	37.84	52.0	2535.0	489.0	1094.0	514.0	3.6591	299200.0	NEAR BAY
7	-122.25	37.84	52.0	3104.0	687.0	1157.0	647.0	3.1200	241400.0	NEAR BAY
8	-122.26	37.84	42.0	2555.0	665.0	1206.0	595.0	2.0804	226700.0	NEAR BAY
9	-122.25	37.84	52.0	3549.0	707.0	1551.0	714.0	3.6912	261100.0	NEAR BAY
10	-122.26	37.85	52.0	2202.0	434.0	910.0	402.0	3.2031	281500.0	NEAR BAY
11	-122.26	37.85	52.0	3503.0	752.0	1504.0	734.0	3.2705	241800.0	NEAR BAY
12	-122.26	37.85	52.0	2491.0	474.0	1098.0	468.0	3.0750	213500.0	NEAR BAY
13	-122.26	37.84	52.0	696.0	191.0	345.0	174.0	2.6736	191300.0	NEAR BAY

图 3: Data Structure

The `info` method is useful to get a quick description of the data, in particular the total number of rows, and each attribute's type and number of non-null values. There are 20640 instances in the dataset, which means that it is fairly small by Machine Learning standards, but it's perfect to get started. Notice that the `total_bedrooms` attribute has only 20433 non-null values, meaning that 207 districts are missing this feature. We will need to take care of this later.

```
In [5]: raw_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
longitude      20640 non-null float64
latitude       20640 non-null float64
housing_median_age  20640 non-null float64
total_rooms    20640 non-null float64
total_bedrooms 20433 non-null float64
population     20640 non-null float64
households     20640 non-null float64
median_income  20640 non-null float64
median_house_value 20640 non-null float64
ocean_proximity 20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

图 4: Data Info

All attributes are numerical, except the `ocean_proximity` field. We can find out what categories exist and how many districts belong to each category by using the `value_counts()` method:

```
In [6]: raw_data['ocean_proximity'].value_counts()

Out[6]: <1H OCEAN      9136
        INLAND       6551
        NEAR OCEAN   2658
        NEAR BAY     2290
        ISLAND        5
        Name: ocean_proximity, dtype: int64
```

图 5: Data Counts

Let's look at the other fields. The describe() method shows a summary of the numerical attributes. (Figure 6). The count, mean, min and max rows are self-explanatory.

```
In [7]: raw_data.describe()

Out[7]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
count	20640.000000	20640.000000	20640.000000	20640.000000	20413.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	206855.816909
std	2.003532	2.135952	12.585558	2191.615252	421.385070	1132.462122	382.329753	1.899822	115395.615874
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	14999.000000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	119600.000000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	179700.000000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	264725.000000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	500001.000000

图 6: Data Describe

Another quick way to get a feel of the type of data you are dealing with is to plot a histogram for each numerical attribute.

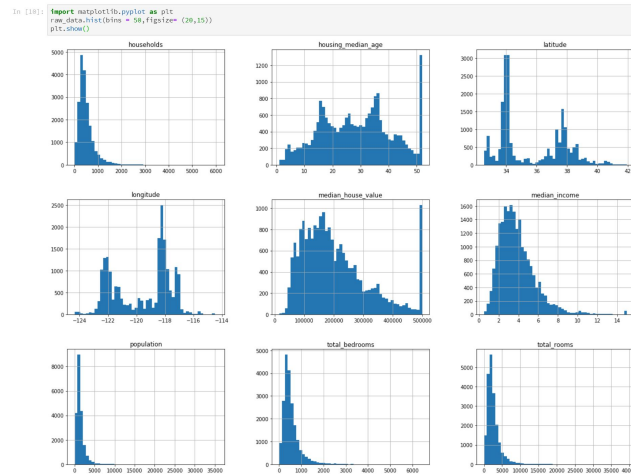


图 7: A histogram for each numerical attribute

4.3 Create a Testset

Create a test set is theoretically quite simple: just pick some instances randomly, typically 20% of the dataset, and set them aside.

```
1 import numpy as np
3 def split_train_test(data, test_ratio):
    shuffled_indices = np.random.permutation(len(data))
5     test_set_size = int(len(data)*test_ratio)
    test_indices = shuffled_indices[:test_set_size]
7     train_indices = shuffled_indices[test_set_size:]
    return data.iloc[train_indices], data.iloc[test_indices]
```

We can then use the function like this:

```
1 train_set, test_set = split_train_test(housing, 0.2)
```

This works, but it's not perfect: if you run the program again, it will generate a different test set. One solution is to save the test set on the first run and then load it in subsequent runs. Another option is to set the random number generator's seed (eg. `np.random.seed(52)`) before calling `np.permutation()`. But both these solutions will break next time you fetch an updated dataset. A common solution is to use each instance's identifier to decide whether or not it should go in the test set. For example, you could compute a hash of each instance's identifier.

```
1 import hashlib
3 def test_set_check(identifier, test_ratio, hash):
    return hash(np.int64(identifier).digest()[-1]) < 256*test_ratio
5
7 def split_train_test_by_id(data, test_ratio, id_column, hash=hashlib.md5):
    ids = data[id_column]
    in_test_set = ids.apply(lambda id_: test_set_check(id_, test_ratio, hash))
9     return data.loc[~in_test_set], data.loc[in_test_set]
```

Unfortunately, the housing dataset does not have an identifier column. The simplest solution is to use the row index as the ID:

```
1 housing_with_id = housing.reset_index()  
3 train_set, test_set = split_train_test_by_id(housing_with_id, 0.2, "index")
```

We can utilize the `train_test_split` function to reach the same effect of the `split` function above.

```
1 from sklearn.model_selection import train_test_split  
2 train_set, test_set = train_test_split(housing, test_size = 0.25, random_state = 0)
```

This is generally fine if your dataset is large enough, but if it is not, you run the risk of introducing a significant sampling bias. For example, when a survey company decides to call 1000 people to ask them a few questions, they don't just pick 1000 people randomly in a phone booth. They try to ensure that these 1000 people are representative of the whole population, e.g. the US population is composed of 51.3% female and 48.7% male, so a survey in the US should try to maintain this ratio in the sample: 513 female and 487 male. This is called *stratified sampling*.

Since the median income is a very important attribute to predict median housing prices. You may want to ensure that the test set is representative of the various categories of incomes in the whole dataset